Triage Report — Quantica

1. Submission Summary

Quantica proposes an AI-powered **private banking pre-review assistant** that continuously monitors client portfolios, detects portfolio drift and life events, and prepares Relationship Manager (RM)-ready rebalancing recommendations with rationale, compliance checks, and auditability before client conversations occur.

- **Industry:** Financial Services
- **Use Case: Portfolio awareness Approach — Event Driven** (private banking portfolio monitoring, actionable-moment detection, and RM task creation)
- **Track: C (inferred)** — Dynamics 365 + Copilot Studio + Power Platform + optional Azure
- **Stage:** Submission Review

Cross-cutting checks

Check	Status	Notes
Environment Track declared	**Warn**	The submission does **not explicitly state the track**, but the stack strongly aligns to **Track C** through **Dynamics 365 CRM, Copilot Studio, Dataverse/Power Platform, and Azure services**. Alignment is good, but the team should explicitly declare Track C.
Industry Use Case aligned	**Pass**	Clear match to **Financial Services, Section 3.9 — Portfolio awareness Approach — Event Driven**. The solution directly targets the private banking pre-review/data-gathering problem using portfolio drift and life-event triggers.
Azure WAF aligned	**Warn**	Good coverage of **Security** and parts of **Operational Excellence** via **Entra ID, Key Vault, Azure Monitor/App Insights**, plus modular/event-driven architecture. However, **Reliability**, **Performance Efficiency**, and especially **Cost Optimisation** are not yet described in enough detail.
Azure WAF AI Pillar aligned	**Warn**	Strong intent on **Agent Design**, **human-in-the-loop**, structured outputs, suitability checks, and auditability. Gaps remain in **model lifecycle**, formal **evaluation**, **grounding provenance**, and **AI observability metrics**.
Power Platform WAF aligned	**Warn**	Use of **Dynamics 365, Dataverse, and Copilot Studio** is appropriate and supports experience optimisation for RMs. However, **Power Platform governance**, **DLP**, **ALM/environment strategy**, and operational controls are not specified.
Microsoft Security & Monitoring	**Warn**	The submission lists **Azure Monitor & App Insights, Key Vault, and Entra ID**, which is a solid baseline. Missing detail on **RBAC boundaries, alerting, privileged access, PII controls, retention, encryption scope, and incident monitoring** across all components.

2. Criteria Scoring

Architecture

Dimension	Score	Justification

Clarity & Flow	**8/10**	The end-to-end flow is easy to follow: portfolio/client signals → life-event detection → analytics/drift detection → recommendation generation → compliance/suitability review → RM action. Agent roles and human approval points are visible, satisfying **Clear Architecture Flow**, though some technical exception paths and state transitions are still high-level.
Platform Fit	**8/10**	The platform choices fit the use case well: **Dynamics 365** for client/life-event context, **Copilot Studio/RM Copilot** for the RM interface, **Azure AI Foundry/Agent Service** for orchestration and reasoning, and Azure data services for storage. This aligns well to **Track C**, but the track is not explicitly declared and the submission is light on how broader integration/governance components are provisioned.
Data / Governance / Security	**7/10**	The submission identifies key stores and controls: **Dataverse, Azure SQL, Blob, ADLS, Entra ID, Key Vault**, audit trail, suitability checks, JSON/schema validation, and approval gates. That is good coverage of **Data & Governance** and **Security**, but an FSI-grade design would also need explicit treatment of **data classification, PII handling, retention, encryption, lineage, access segmentation, and regulatory controls**.
Reliability / Performance / Cost	**6/10**	The design shows promising elements for **Performance** and **Operational Excellence** through event-driven triggers, prioritization, and monitoring with **Azure Monitor/App Insights**. However, the submission does not yet define **failure handling, retry logic, fallback behavior, latency targets, throughput assumptions, token/model routing, or cost controls**, so **Reliability** and **Cost Awareness** remain underdeveloped.
Scalability / Integration / Provisioning	**8/10**	The architecture is modular and explicitly positioned as **globally scalable, locally configured**, which supports **Scalability**. Integrations between CRM, market data, portfolio data, AI agents, and RM-facing tools are defined at a conceptual level, and required services are listed, which helps **Provisioning Readiness**. The main gap is lack of **IaC, deployment topology, and interface/API contract detail**.

Architecture Subtotal: 37/50

Use Case

Dimension	Score	Justification
Value & KPI Impact	**9/10**	This is a strong business problem with clear importance in wealth management. The submission articulates **Problem Significance**, **User Impact**, and **KPI Impact** well: reducing pre-review prep from hours to minutes, scaling RM books, improving proactive service, and maintaining auditability/compliance. It would score even higher with more explicit baseline/target KPI definitions.
Innovation	**8/10**	The solution is more than a generic chatbot: the **dual-trigger model** (market drift + life events), proactive RM tasking, structured recommendation outputs, and embedded compliance workflow create a differentiated approach. It is innovative within the hackathon context, though still closely aligned to the published industry scenario rather than radically novel.

AI Fit	**8/10**	AI is well matched to the parts of the workflow that require reasoning, summarisation, prioritisation, and natural-language interaction, while deterministic thresholds govern drift detection. That is good **AI Fit**. Responsible design is partially addressed through **human-in-the-loop, approval gates, audit trail, structured outputs, and suitability checks**, but the submission lacks explicit detail on **bias, explainability, model evaluation, and safety testing**.
UX Simplicity	**7/10**	The RM-facing experience appears intuitive: dashboard alerts, natural-language copilot interaction, and RM approval before recommendations are surfaced. That supports **User Experience** simplicity. Still, the submission does not yet show the concrete UX flow in enough depth—e.g., card design, approval screens, exception handling, and what the RM sees at each step.
Build–Scale–Reuse	**8/10**	The modular architecture and Microsoft-native stack make the solution plausibly buildable as a hackathon POC, especially if scoped around the highest-value flow. It also has clear **Scalability & Security** potential and reuse value in the form of reusable agents/modules for monitoring, suitability checking, recommendation drafting, and RM workflow support. The main caution is that the full end-to-end scope is ambitious for hackathon time.

Use Case Subtotal: 40/50

Total

Total Score: 77/100

3. Strengths

- Strong alignment to the **Financial Services / Portfolio awareness Approach — Event Driven** scenario, with a problem statement that clearly addresses a real private banking pain point **[Value & KPI Impact; Industry Use Case aligned]**.
- Clear end-to-end business flow from trigger detection through recommendation, compliance review, and RM approval **[Clarity & Flow]**.
- Good Microsoft platform selection: **Dynamics 365 + Copilot Studio + Dataverse + Azure AI Foundry + Azure data services** are well matched to the process and largely consistent with **Track C [Platform Fit]**.
- Compliance is embedded into the flow rather than added later, with suitability checks, audit trail, human approval, and structured outputs **[Data / Governance / Security; AI Fit]**.
- The modular, “start small / scale fast” design supports expansion across markets and asset classes without a full redesign **[Scalability / Integration / Provisioning; Build–Scale–Reuse]**.
- The solution uses AI where it adds value—reasoning, summarisation, recommendation drafting, and RM interaction—while retaining deterministic triggers for drift thresholds **[AI Fit]**.

4. Gaps & Risks

- The submission does not explicitly declare the environment track, even though the architecture appears to fit **Track C**; that ambiguity weakens provisioning readiness **[Platform Fit; Environment Track declared]**.
- For an FSI scenario, data governance detail is still too light: no clear treatment of **PII classification, retention, lineage, encryption scope, access segregation, or compliance mapping [Data / Governance / Security]**.
- Non-functional architecture is under-specified: **retry/fallback logic, SLA/latency targets, failure recovery, throughput assumptions, and AI cost management** are not described **[Reliability / Performance / Cost]**.
- The claim of “hallucination-free recommendations” is stronger than the evidence presented; schema validation and approval gates help, but they do not replace **grounding, evaluation, and observability [AI Fit; Azure WAF AI Pillar aligned]**.
- UX is promising but still conceptual; the submission needs a clearer picture of the RM workflow, approval interaction, and exception handling to fully demonstrate simplicity **[UX Simplicity]**.

- The full scope—portfolio monitoring, life-event detection, analytics, recommendation generation, compliance, and execution support—may be broad for hackathon delivery unless carefully prioritised **[Build–Scale–Reuse]**.
- Despite heavy multi-source data integration and analytics needs, the submission does not explain whether/how **Microsoft Fabric** or an equivalent governed data integration pattern is used **[Platform Fit; Scalability / Integration / Provisioning]**.

5. Triage Verdict

Good — 77/100

6. Top Actionable Improvements

- **Explicitly declare Track C** in the submission and add a one-page architecture/provisioning view showing each component, integration boundary, and deployment dependency, including whether Fabric, API management, or connectors are used **[Platform Fit; Scalability / Integration / Provisioning]**.
- Add an **FSI-grade security and governance matrix** covering client PII, portfolio data, RM actions, and model outputs: RBAC by persona, encryption, retention, lineage, audit scope, and compliance ownership **[Data / Governance / Security; Microsoft Security & Monitoring]**.
- Define key **non-functional targets and controls**: scan frequency, alert latency, retry/fallback behavior, model routing, throughput assumptions, and cost per portfolio/recommendation cycle **[Reliability / Performance / Cost]**.
- Replace the “hallucination-free” statement with **measurable responsible-AI evidence**: grounding sources, evaluation dataset, precision/false-positive targets, approval thresholds, and AI observability dashboards **[AI Fit; Azure WAF AI Pillar aligned]**.
- Narrow the POC to one **high-value demonstrable path**—for example, portfolio drift trigger → recommendation packet → RM approval → audit log—and show the exact RM interaction flow to strengthen feasibility and UX clarity **[UX Simplicity; Build–Scale–Reuse]**.

Important Notice — AI-Assisted Evaluation
This report has been produced using an AI-powered triage system developed by Capgemini as part of the Microsoft Agentic Industry Hackathon 2026 evaluation process. While every effort has been made to ensure consistency and quality, the scores and assessments contained herein are indicative only and should not be treated as definitive or final judgements. AI-generated evaluations may contain inaccuracies, omissions, or contextual misinterpretations.

This report is intended as a supporting tool to assist Capgemini and Microsoft reviewers and also participating teams, judges, and Partner Solution Architects (PSAs) in identifying teams that are currently active and assessing relative submission strength, with the aim of informing — not replacing — human decision-making on stage progression. All final advancement decisions will be made by qualified human reviewers from Capgemini and Microsoft, who will use this report as one input alongside their own professional assessment.

Support & Enquiries: Clemens Reijnen <clemens.reijnen@sogeti.com> | Logan Talbot <logan.talbot@capgemini.com> | v0.0.2

Capgemini